



Relational Database Constraints

Chapter 3 b



Chapter Outline

◆ Relational integrity constraints

- Key constraints
- Entity integrity
- Referential integrity
- Other types of constraints

◆ Example

◆ ER to Relational mapping algorithm

Relational Integrity Constraints



Constraints are *conditions* that must hold on *all* valid relation instances. There are three main types of constraints:

1. **Key** constraints
2. **Entity integrity** constraints
3. **Referential integrity** constraints

Key Constraints

- ◆ Superkey of R: A set of attributes SK of R such that no two tuples *in any valid relation instance* $r(R)$ will have the same value for SK. That is, for any distinct tuples t_1 and t_2 in $r(R)$, $t_1[SK] \neq t_2[SK]$.
- ◆ Key of R: A "minimal" superkey; that is, a superkey K such that removal of any attribute from K results in a set of attributes that is not a superkey.

Key Constraints II

✧ Example: The CAR relation schema:

CAR(Label, Reg#, SerialNo, Make, Model, Year)

has two keys $\text{Key1} = \{\text{Label}, \text{Reg}\# \}$, $\text{Key2} = \{\text{SerialNo}\}$, which are also superkeys. $\{\text{SerialNo}, \text{Make}\}$ is a superkey but *not* a key.

◆ If a relation has *several* **candidate keys**, one is chosen arbitrarily to be the **primary key**. The primary key attributes are *underlined*.

Entity Integrity

- ◆ **Relational Database Schema:** A set S of relation schemas that belong to the same database. S is the *name* of the database. $S = \{R_1, R_2, \dots, R_n\}$
- ◆ **Entity Integrity:** The *primary key attributes* PK of each relation schema R in S cannot have null values in any tuple of $r(R)$. This is because primary key values are used to *identify* the individual tuples.

$$t[\text{PK}] \neq \text{null for any tuple } t \text{ in } r(R)$$

Referential Integrity

- ◆ A constraint involving *two* relations (the previous constraints involve a *single* relation).
- ◆ Used to specify a *relationship* among tuples in two relations: the **referencing relation** and the **referenced relation**.
- ◆ Tuples in the *referencing relation* R_1 have attributes FK (called **foreign key** attributes) that reference the primary key attributes PK of the *referenced relation* R_2 . A tuple t_1 in R_1 is said to **reference** a tuple t_2 in R_2 if $t_1[\text{FK}] = t_2[\text{PK}]$.
- ◆ A referential integrity constraint can be displayed in a relational database schema as a directed arc from $R_1.\text{FK}$ to R_2 .

Referential Integrity II

Statement of the constraint

The value in the foreign key column(s) FK of the **referencing relation** R_1 can be either:

1. a value of an existing primary key value of the corresponding primary key PK in the **referenced relation** R_2 , or.
2. a null.

In case (2), the FK in R_1 should not be a part of its own primary key.

Other Types of Constraints

Semantic Integrity Constraints:

- based on application semantics and cannot be expressed by the model
- E.g., “the max. no. of hours per employee for all projects he on is 56 hrs per week”
- *A constraint specification language* may have to be used to express these
- SQL-99 allows triggers and ASSERTIONS

Example

Employee

FName	MInit	LName	SSN	BDate	Address	Sex	Salary	SuperSSN	DNO
-------	-------	-------	-----	-------	---------	-----	--------	----------	-----

Department

DName	DNumber	MGRSSN	MGRStartDate
-------	---------	--------	--------------

DeptLocations

DNumber	DLocation
---------	-----------

Project

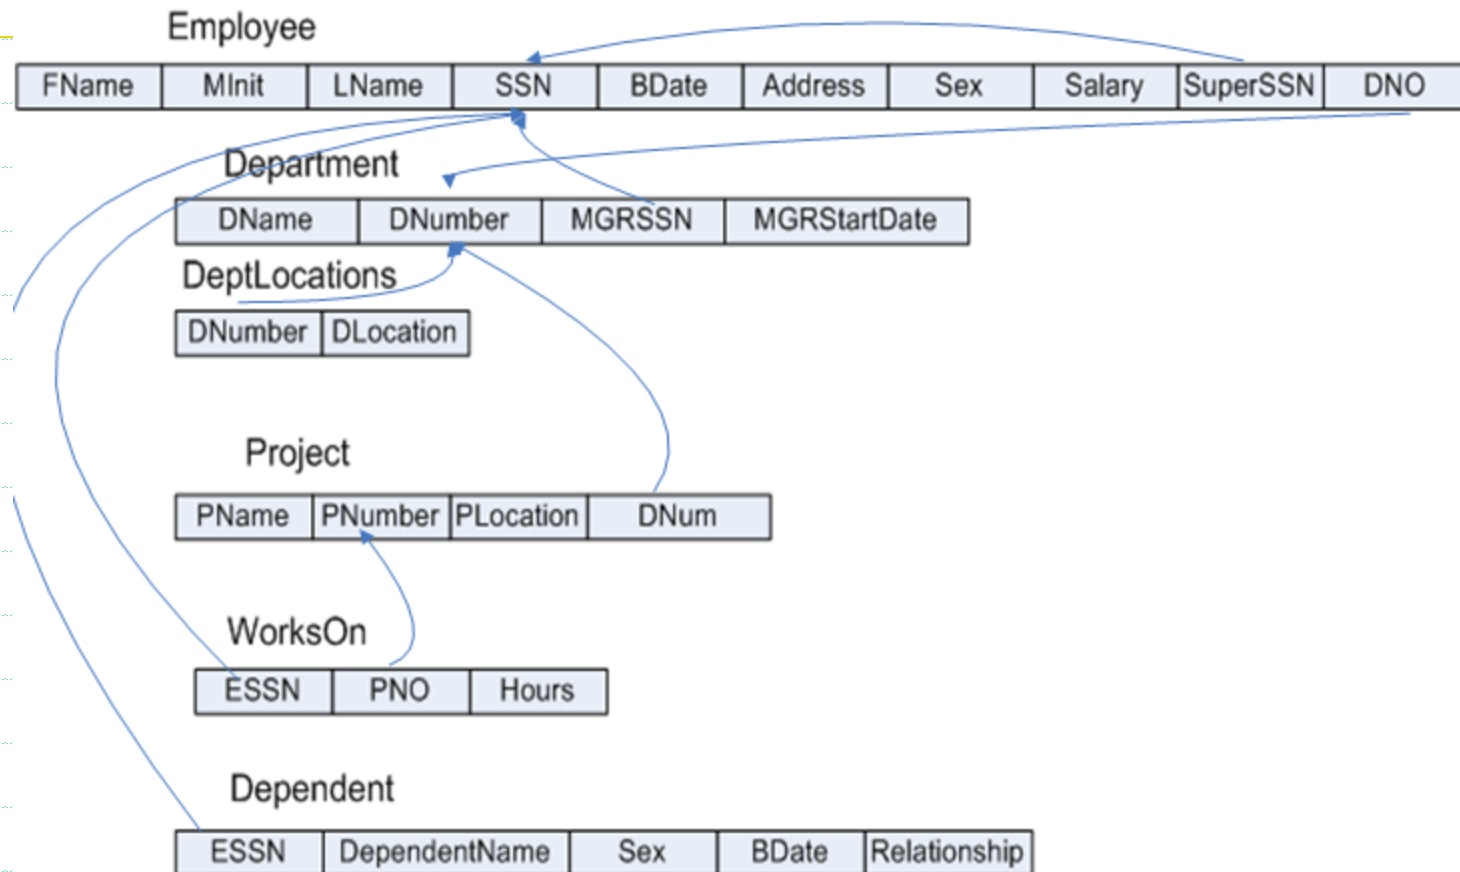
PName	PNumber	PLocation	DNum
-------	---------	-----------	------

WorksOn

ESSN	PNO	Hours
------	-----	-------

Dependent

ESSN	DependentName	Sex	BDate	Relationship
------	---------------	-----	-------	--------------



ER-to-Relational Mapping Algorithm

- **STEP 1:** For each regular (strong) entity type E in the ER schema, create a relation R that includes all the simple attributes of E
- **STEP 2:** For each weak entity W with owner entity E , create a relation R , and include all simple attributes of W as attributes of R .
 - In addition, include as **foreign key** attributes of R the **primary key** attribute(s) of the relation(s) that correspond to the owner entity type(s);

ER-to-Relational Mapping Algorithm

STEP 3: For each binary 1:1 relationship R in the ER schema, identify the relations S and T that correspond to the entity types participating in R.

- Choose one of the relations—S, say—and include as foreign key in S the primary key of T.
- It is better to choose an entity type with total participation in R in the role of S. Include all the simple attributes of the 1:1 relationship type R as attributes of S

ER-to-Relational Mapping Algorithm

- **STEP 4**: For each regular binary 1:N relationship R, identify the relation S that represents the participating entity type at the *N-side* of the relationship.
 - Include as foreign key in S the primary key of the relation T that represents the other entity participating in R.
 - Include any simple attributes of the 1:N relationship type as attributes of S.

ER-to-Relational Mapping Algorithm

- **STEP 5**: For each binary M:N relationship R, create a new relation S to represent R.
 - Include as foreign key attributes in S the primary keys of the relations that represent the participating entity; **their combination will form the primary key of S.**
 - include any simple attributes of the M:N relationship as attributes of S.

ER-to-Relational Mapping Algorithm

- ◆ **STEP 6**: For each **multivalued** attribute A, create a new relation R.
- This relation R will include an attribute corresponding to A,
 - plus the primary key attribute K—as a foreign key in R—of the relation that represents the entity or relationship that has A as an attribute.
 - The primary key of R is the combination of A and K.
If the multivalued attribute is composite, we include its simple components

ER-to-Relational Mapping Algorithm

◆ **STEP 7**: For each n -ary relationship R , where $n > 2$, create a new relation S to represent R .

- Include as foreign key attributes in S the primary keys of the relations that represent the participating entity.
- Also include any simple attributes of the n -ary relationship type as attributes of S .
- The primary key of S is usually a combination of all the foreign keys that reference the relations representing the participating entity types.

ER vs. Relational

ER Model

Entity type

1:1 or 1:N relationship type

M:N relationship type

n-ary relationship type

Simple attribute

Composite attribute

Multivalued attribute

Value set

Key attribute

Relational Model

"Entity" relation

Foreign key (or "relationship" relation)

"Relationship" relation and two foreign keys

"Relationship" relation and n foreign keys

Attribute

Set of simple component attributes

Relation and foreign key

Domain

Primary (or secondary) key